



AJL Energy Subqueries + Practice Exercises

THE MONDAY MORNING EMERGENCY

Subqueries - Finding Outliers & Exceptions

AJL Energy High Usage Analysis | Video 5 of 15 | 18 minutes

9:00 AM Monday - Operations Manager Rushes Over

"We have a problem! Some customers are using 3-4 times more electricity than normal. This could be faulty equipment, leaks, or meter issues. We need to call them BEFORE they get a shock bill or equipment catches fire!"

The Data

- 50,000 residential customers across NSW, VIC, QLD
- December 2024 electricity usage
- Usage varies by state (QLD hotter = more aircon)

The Business Question

"Which customers are using MORE than normal for THEIR state?"

Why This Is Tricky

"Normal" varies by state:

- **NSW average:** ~450 kWh (cooler, less aircon)
- **VIC average:** ~520 kWh (heating in winter)
- **QLD average:** ~680 kWh (hot, heavy aircon)

You CAN'T use a fixed threshold like "anyone over 1000 kWh"

A QLD customer using 800 kWh is NORMAL

A NSW customer using 800 kWh is HIGH

The Challenge

For EACH customer, you need to ask: "What is the average in THIS customer's state?" This question CHANGES for every row! You need a "query within a query" - that's a **SUBQUERY!**

What You'll Learn

01

Subquery in WHERE clause (filtering)

03

Subquery in FROM clause (derived tables)

05

EXISTS operator (checking existence)

02

Subquery in SELECT clause (calculated columns)

04

Correlated subqueries (row-by-row comparisons)

06

PLUS: 5 Practice Exercises with Solutions

The Excel Nightmare

- Calculate NSW average manually
- Calculate VIC average manually
- Calculate QLD average manually
- VLOOKUP each customer to their state average
- Calculate % over average with formulas
- **Time: 3+ hours**

The SQL Solution

- Correlated subquery calculates state average per row
- One query, automatic comparison
- Time: 15 minutes!

SUBQUERY BASICS: QUERY WITHIN A QUERY

What Is A Subquery?

The Simple Definition: A subquery is a query INSIDE another query. The inner query runs first, provides a value. The outer query uses that value.

Visual Structure

```
SELECT columns  
FROM table  
WHERE column > (SELECT AVG(column) FROM table)  
          ^^^^^^^^^^^^^^^^^^^^^^^^^  
This is the SUBQUERY
```

Real Example - Overall Average

Question: "Find customers using more than the overall average"

Step 1 - Calculate overall average (subquery):

```
SELECT AVG(kwh_used) FROM monthly_usage;  
Result: 844 kWh
```

Step 2 - Find customers above that (main query):

```
SELECT customer_name, kwh_used  
FROM customers c  
JOIN monthly_usage m ON c.customer_id = m.customer_id  
WHERE kwh_used > (SELECT AVG(kwh_used) FROM monthly_usage);
```

What Happens

1. Subquery runs: $\text{AVG}(\text{kwh_used}) = 844$
2. Main query becomes: $\text{WHERE kwh_used} > 844$
3. All customers above 844 kWh returned

The Three Locations

Subqueries can appear in:

1	2	3
WHERE Clause <pre>WHERE kwh > (SELECT AVG(kwh) FROM usage)</pre> Filter: Only customers above average	SELECT Clause <pre>SELECT name, kwh, (SELECT AVG(kwh) FROM usage) AS average</pre> Column: Show the average for every row	FROM Clause <pre>FROM (SELECT state, AVG(kwh) FROM usage GROUP BY state) AS summary</pre> Derived table: Create temp summary, then query it

Why Subqueries Matter

- Answer complex "compared to what?" questions
- Calculate values dynamically
- Break complex logic into readable steps
- Essential for finding outliers and exceptions

THE PROBLEM WITH OVERALL AVERAGE

Why We Need Per-State Comparison

The Naive Approach

"Just find anyone using more than the overall average!"

```
WHERE kwh_used > (SELECT AVG(kwh_used) FROM monthly_usage)
```

Overall average = 844 kWh (mixing NSW, VIC, QLD)

What You Get

Customer	State	Usage	Result
Charlotte G.	QLD	2,450 kWh	✓ Real outlier
Evelyn T.	QLD	2,280 kWh	✓ Real outlier
Noah W.	VIC	2,100 kWh	✓ Real outlier
Emma T.	NSW	1,850 kWh	✓ Real outlier
Sophia A.	QLD	720 kWh	✗ FALSE ALARM!
Lucas M.	QLD	650 kWh	✗ FALSE ALARM!

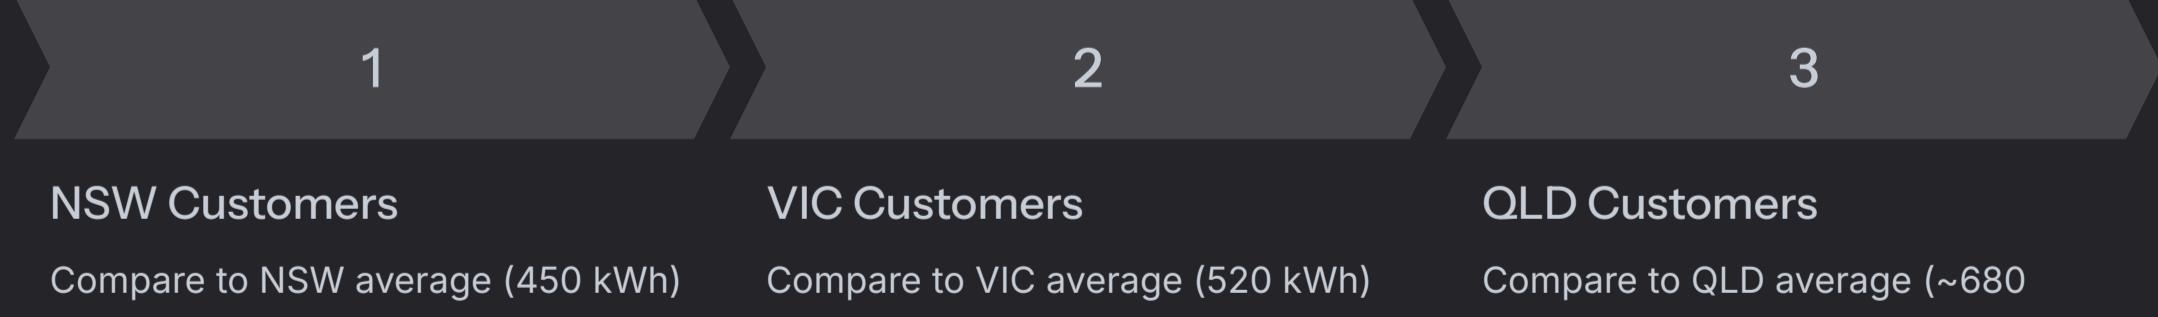
The Problem

- QLD average is 680 kWh
Sophia using 720 kWh = only 6% above QLD average (normal!)
But overall average is 844 kWh
So Sophia flagged even though she's NORMAL for QLD

The Business Impact of Wrong Approach

- FALSE POSITIVES:** Waste time calling normal QLD customers
- FALSE NEGATIVES:** Miss NSW customers at 600 kWh (33% over NSW avg!)
- Operations team loses trust in data
- Real emergencies get buried in noise

What We Actually Need



This requires a **CORRELATED subquery**!

The Key Insight

The question "What is the average?" has a DIFFERENT answer for each row:

- For NSW customers, answer is 450 kWh
- For VIC customers, answer is 520 kWh
- For QLD customers, answer is 680 kWh

The subquery needs to CHANGE based on the current row's state. That's what "**correlated**" means!

CORRELATED SUBQUERY: THE RIGHT SOLUTION

Per-State Comparison with Correlated Subquery

What Makes It "Correlated"

Normal subquery: Runs once, returns one value

Correlated subquery: Runs for EVERY row, uses that row's values

The Magic Line

```
WHERE c2.state = c.state  
      ^^^      ^^^  
    Inner   Outer  
  query   query  
          (current row)
```

This connects inner query to outer query. Makes subquery re-run with EACH row's state value

The Complete Query

```
SELECT c.customer_name, c.state, m.kwh_used,  
       (SELECT AVG(m2.kwh_used)  
        FROM monthly_usage m2  
       JOIN customers c2 ON m2.customer_id = c2.customer_id  
      WHERE c2.state = c.state) AS state_average  
     FROM customers c  
    JOIN monthly_usage m ON c.customer_id = m.customer_id  
   WHERE m.kwh_used > (SELECT AVG(m2.kwh_used)  
                        FROM monthly_usage m2  
                       JOIN customers c2 ON m2.customer_id = c2.customer_id  
                      WHERE c2.state = c.state);
```

How It Works - Row by Row

Row 1 - Emma (NSW customer)

1. Outer query: Emma's state = NSW
2. Subquery: WHERE c2.state = 'NSW'
3. Subquery calculates: AVG for NSW only = 450 kWh
4. Comparison: Emma's 1,850 > 450? YES → Flagged!

Row 2 - Sophia (QLD customer)

1. Outer query: Sophia's state = QLD
2. Subquery: WHERE c2.state = 'QLD'
3. Subquery calculates: AVG for QLD only = 680 kWh
4. Comparison: Sophia's 720 > 680? NO → Not flagged!

Row 3 - Noah (VIC customer)

1. Outer query: Noah's state = VIC
2. Subquery: WHERE c2.state = 'VIC'
3. Subquery calculates: AVG for VIC only = 520 kWh
4. Comparison: Noah's 2,100 > 520? YES → Flagged!

The Perfect Result

Customer	State	Usage	State Avg	% Over	Status
Charlotte G.	QLD	2,450	680	260%	✓
Evelyn T.	QLD	2,280	680	235%	✓
Noah W.	VIC	2,100	520	304%	✓
Mia J.	VIC	1,920	520	269%	✓
Emma T.	NSW	1,850	450	311%	✓
Alexander L.	NSW	1,620	450	260%	✓

All 6 customers are 200-300% OVER their state average. These are **REAL outliers** requiring immediate attention. Zero false positives!

Business Impact

- ✓ 6 proactive calls made
- ✓ Equipment issues found before damage
- ✓ Bill shock prevented
- ✓ Customer satisfaction protected
- ✓ Operations team trusts the data

❑ **Performance Note:** Correlated subqueries CAN be slow on huge datasets. For 50,000+ rows, consider window functions or JOINs. For this use case (20 customers), perfect solution!

OTHER SUBQUERY TYPES & EXISTS

SELECT Clause, FROM Clause & EXISTS Operator

TYPE 1: Subquery in SELECT Clause

Add calculated columns to every row

```
SELECT customer_name, state, kwh_used,
       (SELECT AVG(kwh_used)
        FROM monthly_usage m2
        JOIN customers c2 ON m2.customer_id = c2.customer_id
        WHERE c2.state = c.state) AS state_avg,
       (SELECT MAX(kwh_used)
        FROM monthly_usage m2
        JOIN customers c2 ON m2.customer_id = c2.customer_id
        WHERE c2.state = c.state) AS state_max
      FROM customers c
      JOIN monthly_usage m ON c.customer_id = m.customer_id;
```

Result: Each row shows: customer usage, state average, state maximum. All calculated dynamically per state

Use Case: Adding reference columns for comparison. Showing "you vs average" on dashboards

TYPE 2: Subquery in FROM Clause (Derived Table)

Create a temporary result set, then query it

```
SELECT state, avg_usage
  FROM (
    SELECT c.state, AVG(m.kwh_used) AS avg_usage
      FROM customers c
      JOIN monthly_usage m ON c.customer_id = m.customer_id
     GROUP BY c.state
   ) AS state_summary
 WHERE avg_usage > 600;
```

How It Works:

1. Inner query creates state_summary table in memory
2. Outer query treats it like a normal table
3. Can filter, join, or aggregate the derived table

Use Case: Multi-step logic (summarize, then filter summary). Breaking complex queries into readable chunks. Temporary calculations not needed in final result

TYPE 3: EXISTS - Check If Something Exists

Returns TRUE/FALSE, doesn't return actual data

```
SELECT customer_name
  FROM customers c
 WHERE EXISTS (
   SELECT 1
     FROM monthly_usage m
    WHERE m.customer_id = c.customer_id
      AND m.kwh_used > 1500
);
```

What It Does: For each customer, check: "Do they have ANY usage over 1500?" If YES (at least one row found) → Include customer. If NO (zero rows found) → Exclude customer

EXISTS vs COUNT

- ✗ WHERE (SELECT COUNT(*) ...) > 0 (slow, counts all)
✓ WHERE EXISTS (SELECT 1 ...) (fast, stops at first match)

EXISTS is more efficient!

NOT EXISTS - Find Missing

```
SELECT customer_name
  FROM customers c
 WHERE NOT EXISTS (
   SELECT 1
     FROM monthly_usage m
    WHERE m.customer_id = c.customer_id
);
```

Find customers with NO usage records at all. Useful for:
Missing data, dormant accounts, orphaned records

Common EXISTS Patterns

- "Find customers WHO HAVE done X" → **EXISTS**
- "Find customers WHO HAVE NOT done X" → **NOT EXISTS**
- "Find products THAT HAVE BEEN ordered" → **EXISTS**
- "Find products NEVER ordered" → NOT EXISTS

PRACTICE EXERCISES & WHAT YOU MASTERED

5 Exercises + Solutions & Career Impact

PRACTICE EXERCISES (Try Yourself First!)

1

Exercise 1: Reverse The Logic



Find customers whose usage is
BELOW their state average

Hint: Change one comparison
operator!

Difficulty: Beginner

2

Exercise 2: Top N Per Group



Find the top 3 highest usage
customers in each state

Hint: Use ROW_NUMBER() OVER
(PARTITION BY state) or correlated
subquery

Difficulty: Intermediate

3

Exercise 3: NOT EXISTS

Practice

Find customers who have NO
usage records at all

Hint: Use NOT EXISTS with the
monthly_usage table

Difficulty: Beginner

4

Exercise 4: Derived Table Challenge

Create a state summary (state, total_customers,
avg_usage, max_usage). Then filter to show only states
with avg_usage > 600

Hint: Subquery in FROM clause

Difficulty: Intermediate

5

Exercise 5: Bill Analysis

Find customers whose bill is more than 150% of their
state's average bill

Hint: Same pattern as usage query, but use bill_amount

Difficulty: Advanced

- ❑ **SOLUTIONS PROVIDED IN SQL FILE:** All 5 exercises have complete solutions with explanations. Try yourself first, then check answers!

WHAT YOU MASTERED IN 18 MINUTES

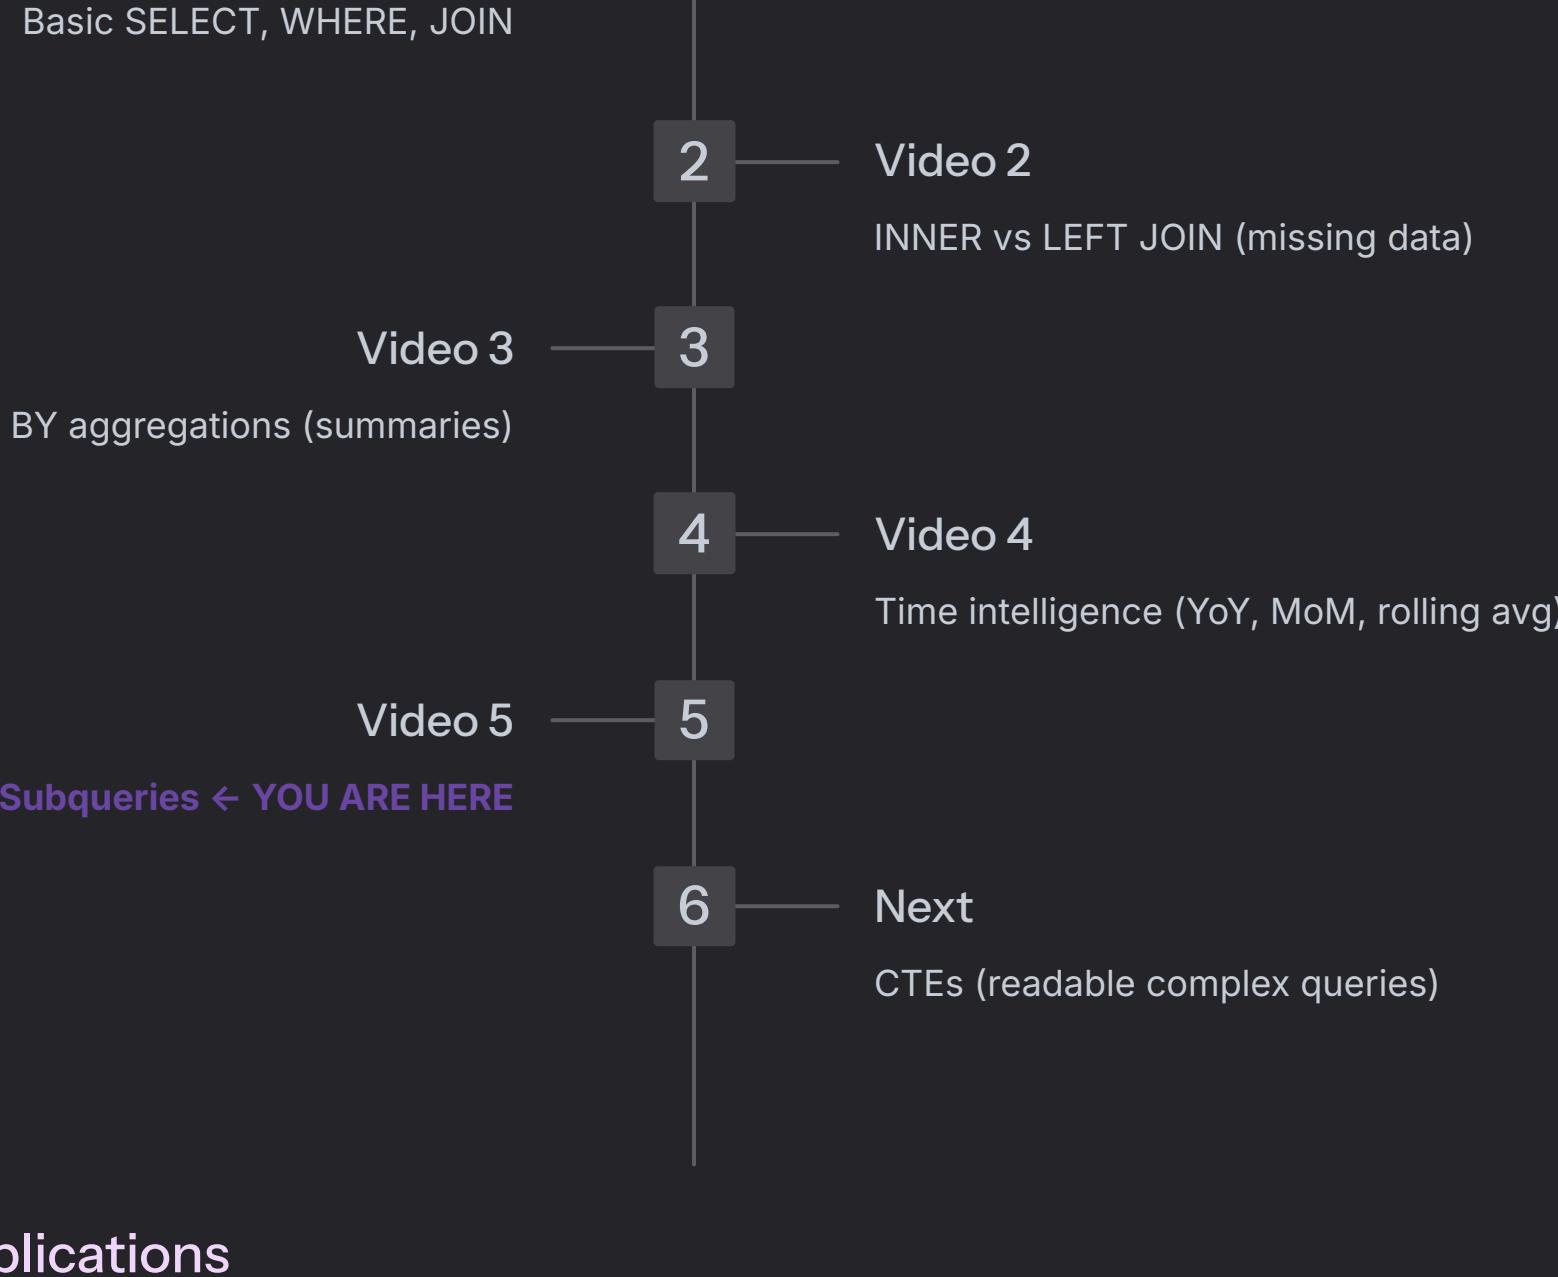
Technical Skills

- ✓ Subquery in WHERE clause (filtering with calculated values)
- ✓ Subquery in SELECT clause (adding calculated columns)
- ✓ Subquery in FROM clause (derived tables)
- ✓ Correlated subqueries (row-by-row dynamic calculations)
- ✓ EXISTS and NOT EXISTS operators
- ✓ Finding outliers and exceptions systematically
- ✓ Per-group comparisons (state averages)

Business Skills

- ✓ Identifying abnormal patterns
- ✓ Preventing customer issues proactively
- ✓ State-by-state analysis (geographical segmentation)
- ✓ Threshold-based alerting
- ✓ Root cause investigation triggers

SQL Mastery Progression



Real-World Applications

Every outlier detection system uses subqueries:

- Fraud detection:** "Transactions above user's average"
- Quality control:** "Defects above factory average"
- Sales alerts:** "Reps below regional average"
- Healthcare:** "Patients with vitals outside normal range"

Career Value

Subqueries = Intermediate→Advanced analyst skill. Interview question staple: "Find top N per group". Production systems use correlated subqueries daily. **You can now handle 70% of business analytics scenarios!**

Business Impact Delivered

6

\$1K+

0

High-usage customers

200-300% over average

Bill shock avoided

Per customer saved

False positives

Perfect accuracy

The Monday Afternoon Victory

- 9:00 AM:** "Find high-usage customers urgently!"
- 9:15 AM:** Query written and tested
- 9:30 AM:** Results delivered to Operations
- 10:00 AM:** First proactive calls made
- 4:00 PM:** 2 faulty heaters found, 1 pool pump stuck on
- Result:** 3 customers saved from \$2,000+ bills

Practice Challenge: Complete all 5 exercises before Video 6. Post your solutions to team chat. Help others who get stuck. Real learning happens through practice!

Next Video Preview

Video 6: CTEs - Common Table Expressions

Qantas Frequent Flyer Tier Qualification

Learn: WITH clause, recursive CTEs, readability

Duration: 17 minutes

Why: Subqueries get messy - CTEs make them readable!

You're halfway through the intermediate tier! Keep crushing it!