

DOCKER

CSE5006 – LAB 2

TABLE OF CONTENTS

1. What is Docker?	3
2. Running a “Hello World” Container	3
2.1. Exercise 1.....	6
3. Creating a Fortune Cookie Image	6
3.1. Humble Beginnings	6
3.2. A Quick Aside: Networking Essentials	10
3.2.1. Exercise 2.....	12
3.3. Running a Simple Node.js Application	12
3.3.1. Exercise 3.....	17

3.4. Adding Fortune Cookie Messages	17
3.4.1. Exercise 4.....	18
3.5. Using Docker Compose	18
3.6. Sharing Directories via Volumes.....	19
3.6.1. Exercise 5.....	21

1. WHAT IS DOCKER?

Software environments can vary greatly, with countless combinations of operating systems, libraries, executables, and configurations. This can cause issues where locally developed applications behave differently in production. Docker bundles applications with their software environment into a self-contained image to address this problem.

2. RUNNING A “HELLO WORLD” CONTAINER

When purchasing a computer without pre-installed software, the typical first step is to install an operating system from a disc or USB, such as Windows 7.

In Docker, an image is like a Windows 7 DVD. Docker images are based on Linux and are lightweight (usually a few hundred megabytes). You can create a container from an image using the Docker command-line tool. Docker allows you to create multiple containers with isolated environments on one computer. For example, you can have containers for compiling Java programs, running scientific simulations, and playing NetHack. If you delete files in the NetHack container, the Java container remains unaffected. Later, you can create a new NetHack container from the original image.

Think of Docker images and containers as classes and objects in object-oriented programming. Images are like blueprints for creating classes, while containers are instances of those classes.

Docker Hub is an online registry hosting thousands of free Docker images that can be downloaded and used to create containers. Let's take a look at how we would go about doing exactly that.

Firstly, we will tell Docker to download the hello-world image. You can read about what's in this particular image by visiting https://hub.docker.com/_/hello-world/.

```
docker pull hello-world
```

```
vboxuser@CSE5006:~$ docker pull hello-world
Using default tag: latest
latest: Pulling from library/hello-world
Digest: sha256:dcba6daec718f547568c562956fa47e1b03673dd010fe6ee58ca806767031d1c
Status: Image is up to date for hello-world:latest
docker.io/library/hello-world:latest
vboxuser@CSE5006:~$
```

After a few moments, you should be informed that the image has been downloaded. To verify if the image has been successfully downloaded, you can use the following script

```
vboxuser@CSE5006:~$ docker image ls
REPOSITORY    TAG       IMAGE ID       CREATED        SIZE
hello-world   latest    9c7a54a9a43c   3 months ago   13.3kB
vboxuser@CSE5006:~$
```

Now we can create a new container from the image.

```
docker create --name=my-container hello-world
```

```
vboxuser@CSE5006:~$ docker create --name=my-container hello-world
5f96e5e90a7567d45d575beb05f015751c80d526d3cf3ac808851fecc79e65ec
```

Note: Do not add spaces when assigning value to a parameter to prevent the following error.

```
vboxuser@CSE5006:~$ docker create --name = my-container hello-world
Unable to find image 'my-container:latest' locally
Error response from daemon: pull access denied for my-container, repository does not exist or may require 'docker log
in': denied: requested access to the resource is denied
vboxuser@CSE5006:~$
```

To check that the container was created successfully, issue the following command.

```
docker ps --all
```

This will list all Docker containers on the system, including information like ID, status, and the image from which they were created. This command is incredibly useful, and I recommend running it often to check up on your containers.

You can use either `$ docker ps` or `$ docker container ls` to show the running containers in your system. Adding `-all` parameter will show all existing containers (running or stop).

```
vboxuser@CSE5006:~$ docker container ls
CONTAINER ID   IMAGE     COMMAND   CREATED   STATUS    PORTS   NAMES
vboxuser@CSE5006:~$ docker container ls -a
CONTAINER ID   IMAGE     COMMAND   CREATED      STATUS      PORTS      NAMES
5f96e5e90a75   hello-world   "/hello"   2 minutes ago   Created           my-container
vboxuser@CSE5006:~$ docker ps -a
CONTAINER ID   IMAGE     COMMAND   CREATED      STATUS      PORTS      NAMES
5f96e5e90a75   hello-world   "/hello"   2 minutes ago   Created           my-container
vboxuser@CSE5006:~$ docker ps -all
CONTAINER ID   IMAGE     COMMAND   CREATED      STATUS      PORTS      NAMES
5f96e5e90a75   hello-world   "/hello"   2 minutes ago   Created           my-container
```

Now it's time to run the main program inside the container.

```
docker start --attach my-container
```

```
vboxuser@CSE5006:~$ docker start --attach my-container

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
   (amd64)
3. The Docker daemon created a new container from that image which runs the
   executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it
   to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/
```

If all went well, you should be presented with a "Hello from Docker" message accompanied by some other informative text. Note that although the program has terminated, the container you created still exists. Using a Docker command you have already learned (hint: it rhymes with "locker tree-mess hall"), check the status of the container. Does the status make sense given what we have done?

To actually remove the container from the system, another command is required.

```
docker rm my-container
```

```
vboxuser@CSE5006:~$ docker rm my-container
my-container
vboxuser@CSE5006:~$ docker ps -a
CONTAINER ID   IMAGE     COMMAND   CREATED   STATUS    PORTS   NAMES
vboxuser@CSE5006:~$
```

Confirm that

the container is no longer present by listing Docker containers. For convenience, Docker provides a handy shortcut command which pulls the image if it doesn't exist, creates a container, runs it, and removes it upon termination. This allows us to achieve what previously took us four commands (pull, create, start, rm) with only one.

```
docker run --rm --name=my-container hello-world
```

You should see the same message as before.

```
vboxuser@CSE5006:~$ docker run --rm --name=my-container hello-world

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
   (amd64)
3. The Docker daemon created a new container from that image which runs the
   executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it
   to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/

vboxuser@CSE5006:~$
```

2.1. EXERCISE 1

Create (but don't run) a hello-world container without specifying a name (omit the `--name` option completely). What is reported in the name field when you list Docker containers? Try removing the container using its ID rather than its name. Does it work? Put your answers from this and all other exercises in a text file called EXERCISES and make sure that it's committed to your repository.

3. CREATING A FORTUNE COOKIE IMAGE

For the remainder of this lab, we will be creating a simple web application. Based on what you learned in the first lab, use Git to keep track of changes and store your code on GitHub. Your demonstrator will be checking your use of version control, so make commits at your discretion as you work through the lab.

3.1. HUMBLE BEGINNINGS


Using images from Docker Hub is great when you want to run pre-existing software. However, when it comes to running your own custom programs, Docker provides a solution. Similar to object-oriented programming, Docker allows you to extend existing images, inheriting all the components present in the base image. This way, you can customize the image by adding whatever you need, including dependencies and your own applications.

Today, our goal is to create a very simple website that displays a new fortune cookie message every time the page is loaded, and then build it into a Docker image. To start, we need to establish a basic project structure.


1. Create a new directory in Lab02 called fortune-cookie, and initialize a Git repository.

```
vboxuser@CSE5006:~/Documents$ mkdir -p ~/Documents/Lab02/fortune-cookie
vboxuser@CSE5006:~/Documents$ cd ~/Documents/Lab02/fortune-cookie/
vboxuser@CSE5006:~/Documents/Lab02/fortune-cookie$ git init
hint: Using 'master' as the name for the initial branch. This default branch name
hint: is subject to change. To configure the initial branch name to use in all
hint: of your new repositories, which will suppress this warning, call:
hint:     git config --global init.defaultBranch <name>
hint: Names commonly chosen instead of 'master' are 'main', 'trunk' and
hint: 'development'. The just-created branch can be renamed via this command:
hint:     git branch -m <name>
Initialized empty Git repository in /home/vboxuser/Documents/Lab02/fortune-cookie/.git/
vboxuser@CSE5006:~/Documents/Lab02/fortune-cookie$
```

2. Start Visual Studio Code and open the newly created folder.



3. From inside Visual Studio Code, create three new files called **Dockerfile**, **package.json**, and **server.js**.



Put the following script in Dockerfile, and leave the rest of the files blank

```
# Base this image on an official Node.js long term support image.
FROM node:18.16.0-alpine

# Install Tini using Alpine Linux's package manager, apk.
RUN apk add --no-cache tini

# Use Tini as the init process. Tini will take care of system stuff
# for us, like forwarding signals and reaping zombie processes.
ENTRYPOINT ["/sbin/tini", "--"]
```

The first part of the file tells Docker that we want our image to be based on **node:18.0.0-alpine**, which is an official Alpine Linux-based image containing Node.js. More on Node.js later; for now, suffice to say that we will be using it to make our web application.

The rest of the file tells Docker how we want to modify the base image. Here we are downloading a small program called Tini, which will act as the init process. It is a good idea to use a proper init process with Docker, but we won't go into the details here. However, you should pay attention to how we add Tini to the Docker image. The RUN step runs a command as if we had a terminal inside the Docker image - in this case, we are simply downloading and installing Tini with Alpine Linux's package manager, **apk**. The **ENTRYPOINT** step tells Docker that when we run a container, we want it to use Tini as the init process.

Start a terminal in the fortune-cookie directory and build our custom Docker image.

```
docker build --tag=fortune-cookie .
```


You should see each step run in order, ending with a message that indicates the image was successfully built. The tag option gives our image a name, just like the hello-world image we encountered earlier.

Notes: Make sure you save the previous changes in Dockerfile before running the docker application to avoid getting the following error:

```
vboxuser@CSE5006:~/Documents/Lab02/fortune-cookie$ docker build --tag=fortune-cookie .
[+] Building 0.1s (2/2) FINISHED                                docker:default
=> [internal] load build definition from Dockerfile              0.0s
=> => transferring dockerfile: 31B                               0.0s
=> [internal] load .dockerignore                                0.0s
=> => transferring context: 2B                                    0.0s
ERROR: failed to solve: the Dockerfile cannot be empty
```

We can confirm that the image is now available on the system by listing all Docker images.

```
docker images
```

```
vboxuser@CSE5006:~/Documents/Lab02/fortune-cookie$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
fortune-cookie	latest	3354c7691bb9	About a minute ago	175MB
hello-world	latest	9c7a54a9a43c	3 months ago	13.3kB

To test the image, we can run it by telling it to start a shell session inside a container (sh is a shell that's already installed in the base image).

```
docker run --rm -it fortune-cookie sh
```

You will be presented with a shell prompt. You can verify that you are indeed inside the image by checking that Tini is available.

```
tini -h
```

If you see usage instructions for Tini, then everything is going according to plan.

```

vboxuser@CSE5006:~/Documents/Lab02/fortune-cookie$ docker run --rm -it fortune-cookie sh
/ # tini -h
tini (tini version 0.19.0)
Usage: tini [OPTIONS] PROGRAM -- [ARGS] | --version

Execute a program under the supervision of a valid init process (tini)

Command line options:

--version: Show version and exit.
-h: Show this help message and exit.
-s: Register as a process subreaper (requires Linux >= 3.4).
-p SIGNAL: Trigger SIGNAL when parent dies, e.g. "-p SIGKILL".
-v: Generate more verbose output. Repeat up to 3 times.
-w: Print a warning when processes are getting reaped.
-g: Send signals to the child's process group.
-e EXIT_CODE: Remap EXIT_CODE (from 0 to 255) to 0.
-l: Show license and exit.

Environment variables:

TINI_SUBREAPER: Register as a process subreaper (requires Linux >= 3.4).
TINI_VERBOSITY: Set the verbosity level (default: 1).
TINI_KILL_PROCESS_GROUP: Send signals to the child's process group.

/ # █

```

Our image (and therefore the container) is still quite bare at the moment. Press **Ctrl+D** to exit the shell so we can continue our work.

3.2. A QUICK ASIDE: NETWORKING ESSENTIALS

Computers are connected over a network using sockets. This is also true for connecting to web servers, as we will demonstrate in this section.

If you provide me with a URL, such as <http://itcorp.com:80/index.html>, for instance, I can already navigate to the associated website using a web browser. However, what may be less familiar to you is how this process involves low-level socket communication to retrieve the relevant content. First, let's break down the anatomy of a URL.

<http://itcorp.com:80/index.html>
 ↑ ↑ ↑ ↑
 Scheme Hostname Port Path

Part	Description
<i>Scheme</i>	The protocol to use when communicating with the server
<i>Hostname</i>	The name of the server we wish to connect to.
<i>Port</i>	Port number on the server to use for communication.
<i>Path</i>	The resource we are requesting from the server.

One part of URLs that you might not encounter frequently is the port number. This is because when an HTTP URL doesn't explicitly include a port, the default value of port 80 is assumed. So, <http://itcorp.com:80/index.html> and <http://itcorp.com/index.html> are equivalent. Port numbers enable multiple server applications to operate on the same host simultaneously, as each application can be assigned a different port.

Although the hostname is a convenient and human-readable way of expressing the name of a server, computers actually communicate using IP addresses. To obtain an IP address from a domain name, we need to query a Domain Name System (DNS) that will map the hostname to the corresponding IP address. We can use the `nslookup` command to perform this lookup explicitly.

```
nslookup itcorp.com
```

```
choiruzain@MacMarichoy-7 fortune-cookie % nslookup itcorp.com
Server:          10.128.8.21
Address:         10.128.8.21#53
```

```
Non-authoritative answer:
Name:   itcorp.com
Address: 134.173.42.59
```

Under the "Non-authoritative answer" section of the output, you should see the IP address associated with itcorp.com. Now we have the two essential pieces of information required to initialize socket communication with the server: the IP address and the port number (default HTTP port 80) of the web application.

```
telnet <ip_address> 80
```

The telnet command utilizes sockets to connect to the application running on port 80 of the itcorp.com server. Please type in the following HTTP request, which includes the path of the desired resource. Our web browsers handle this process automatically when we visit a website.

```
choiruzain@MacMarichoy-7 fortune-cookie % telnet 134.173.42.59 80
Trying 134.173.42.59...
Connected to mallet.cs.hmc.edu.
Escape character is '^]'.
_
```

```
GET /index.html HTTP/1.0
```

Hit enter twice. You should be presented with the same HTML response you received when entering the URL in a web browser.

```

choiruzain@MacMarichoy-7 fortune-cookie % telnet 134.173.42.59 80
Trying 134.173.42.59...
Connected to mallet.cs.hmc.edu.
Escape character is '^]'.
GET /index.html HTTP/1.0

HTTP/1.1 200 OK
Date: Mon, 12 Aug 2024 05:54:48 GMT
Server: Apache/2.4.51 (Linux/SUSE)
Last-Modified: Sat, 28 Jul 2012 01:35:16 GMT
ETag: "551-4c5d9d7351d00"
Accept-Ranges: bytes
Content-Length: 1361
Connection: close
Content-Type: text/html

<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01//EN"
"http://www.w3.org/TR/html4/strict.dtd">
<HTML LANG="en-US">
<HEAD>
  <META HTTP-EQUIV="Content-Type" CONTENT="text/html; charset=ISO-8859-1">
  <META NAME="keywords" CONTENT="bicycle gear design, freewheel,
  cassette, chainwheel">
  <META NAME="description" CONTENT="Javascript-based bicycle gear design">
  <STYLE>
    <!--
    -->
  </STYLE>
  <TITLE>
    Bicycle Gear Design
  </TITLE>
</HEAD>

<BODY>
<H1>Bicycle Gear Design</H1>
<P>
Bestcogs.com is your one-stop site for designing bicycle gearing systems.
We aim to offer a simple, flexible interface that will let you design
an optimal gear setup for your bike.

<H2>Get Started</H2>

<P>
If you just want to design a gearing system,
<A HREF="gears.html">click here</A> and choose your chainwheels,
cassette/freewheel parameters, and gearing preferences.
Your ideal freewheel options will be calculated for you.

<P>
Or perhaps you want <A HREF="instructions.html">more detailed
instructions.</A> That's often the best way to approach things
(RTFM, as they say).

<H2>Other Information</H2>

<P>
This site is still under development. But we already have some
<A HREF="faqs.html">FAQs</A> about the site and about bicycle gears.
We also debunk some <A HREF="myths.html">common myths and
misconceptions</A> about gearing.

</BODY>
</HTML>
Connection closed by foreign host.

```

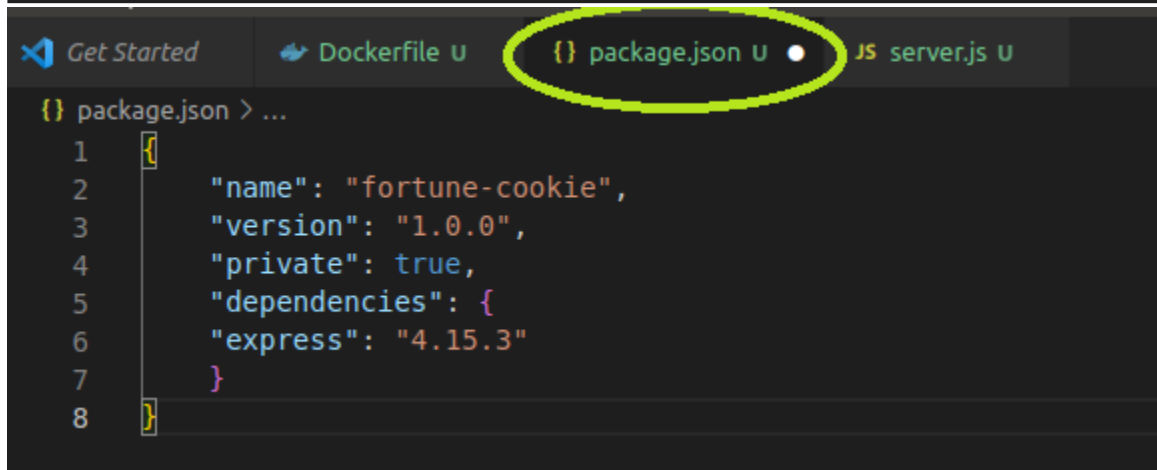
3.2.1. EXERCISE 2

Try finding "the" IP address for amazon.com with **nslookup**. What do you find, and why do you think that might be the case for a big company like Amazon? Don't forget to note your answers.

3.3. RUNNING A SIMPLE NODE.JS APPLICATION

Node.js is a runtime environment for developing JavaScript applications that run on servers. Dependencies for a Node application are specified in the **package.json** file, which we will create now. Since this is a simple project, we only have one dependency - Express, which is a minimalist web framework. While we're at it, we can also add a project name and version number and indicate that the project is not intended for public release.

```
{
  "name": "fortune-cookie",
  "version": "1.0.0",
  "private": true,
  "dependencies": {
    "express": "4.15.3"
  }
}
```



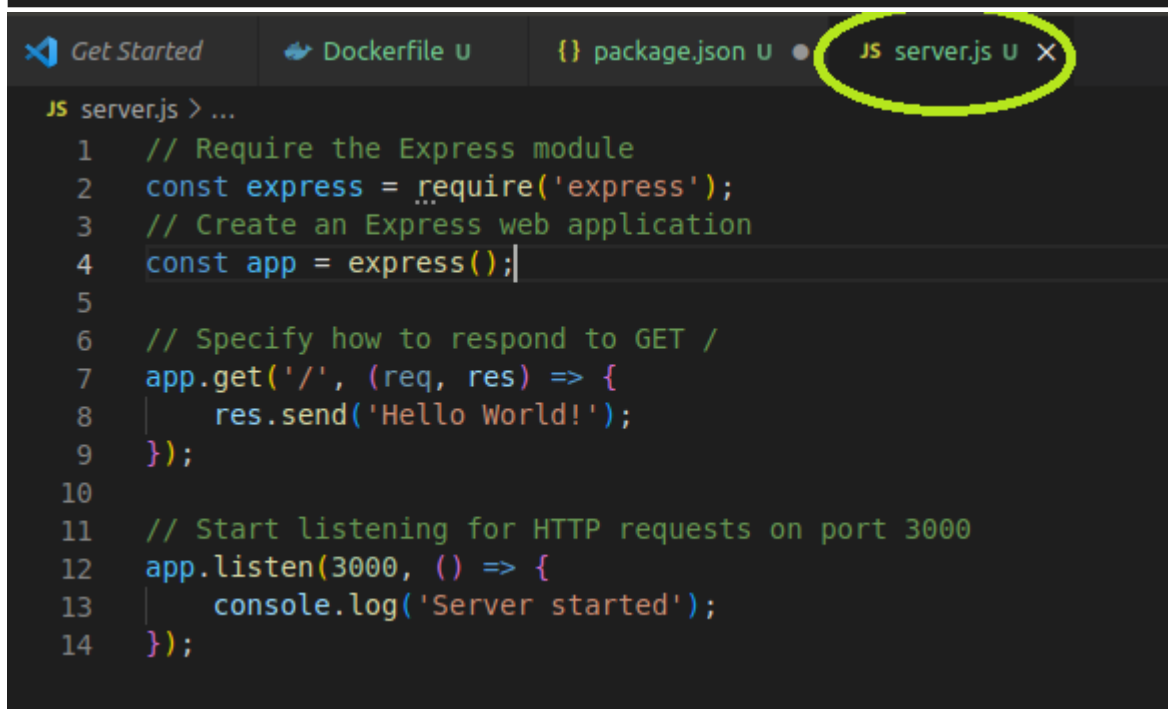
```
{ } package.json > ...
1  {
2      "name": "fortune-cookie",
3      "version": "1.0.0",
4      "private": true,
5      "dependencies": {
6          "express": "4.15.3"
7      }
8  }
```

With that over with, we can create a simple web application just to check that Express is working. Create **server.js** and add a basic "Hello World" server program.

```
// Require the Express module
const express = require('express');
// Create an Express web application
const app = express();

// Specify how to respond to GET /
app.get('/', (req, res) => {
  res.send('Hello World!');
});

// Start listening for HTTP requests on port 3000
app.listen(3000, () => {
  console.log('Server started');
});
```



```
JS server.js > ...
1  // Require the Express module
2  const express = require('express');
3  // Create an Express web application
4  const app = express();
5
6  // Specify how to respond to GET /
7  app.get('/', (req, res) => {
8    res.send('Hello World!');
9  });
10
11 // Start listening for HTTP requests on port 3000
12 app.listen(3000, () => {
13   console.log('Server started');
14 });
```

Let's take a moment to break this down. Firstly, we require the Express module and use it to create an app object. Secondly, we specify that when the server receives a request on the index of the site, it should respond with the message "Hello World!". Lastly, we start the server and listen for requests on port 3000 (as opposed to the default HTTP port 80).

Now that we actually have some custom software, we need to add it to our Docker image. Open up the Dockerfile again and add the following lines to the bottom, below the steps already there.

```
# ...steps from before...

# Create a working directory for our application.
RUN mkdir -p /app
WORKDIR /app

# Install the project's NPM dependencies.
COPY package.json /app/
RUN npm --silent install
RUN mkdir /deps && mv node_modules /deps/node_modules

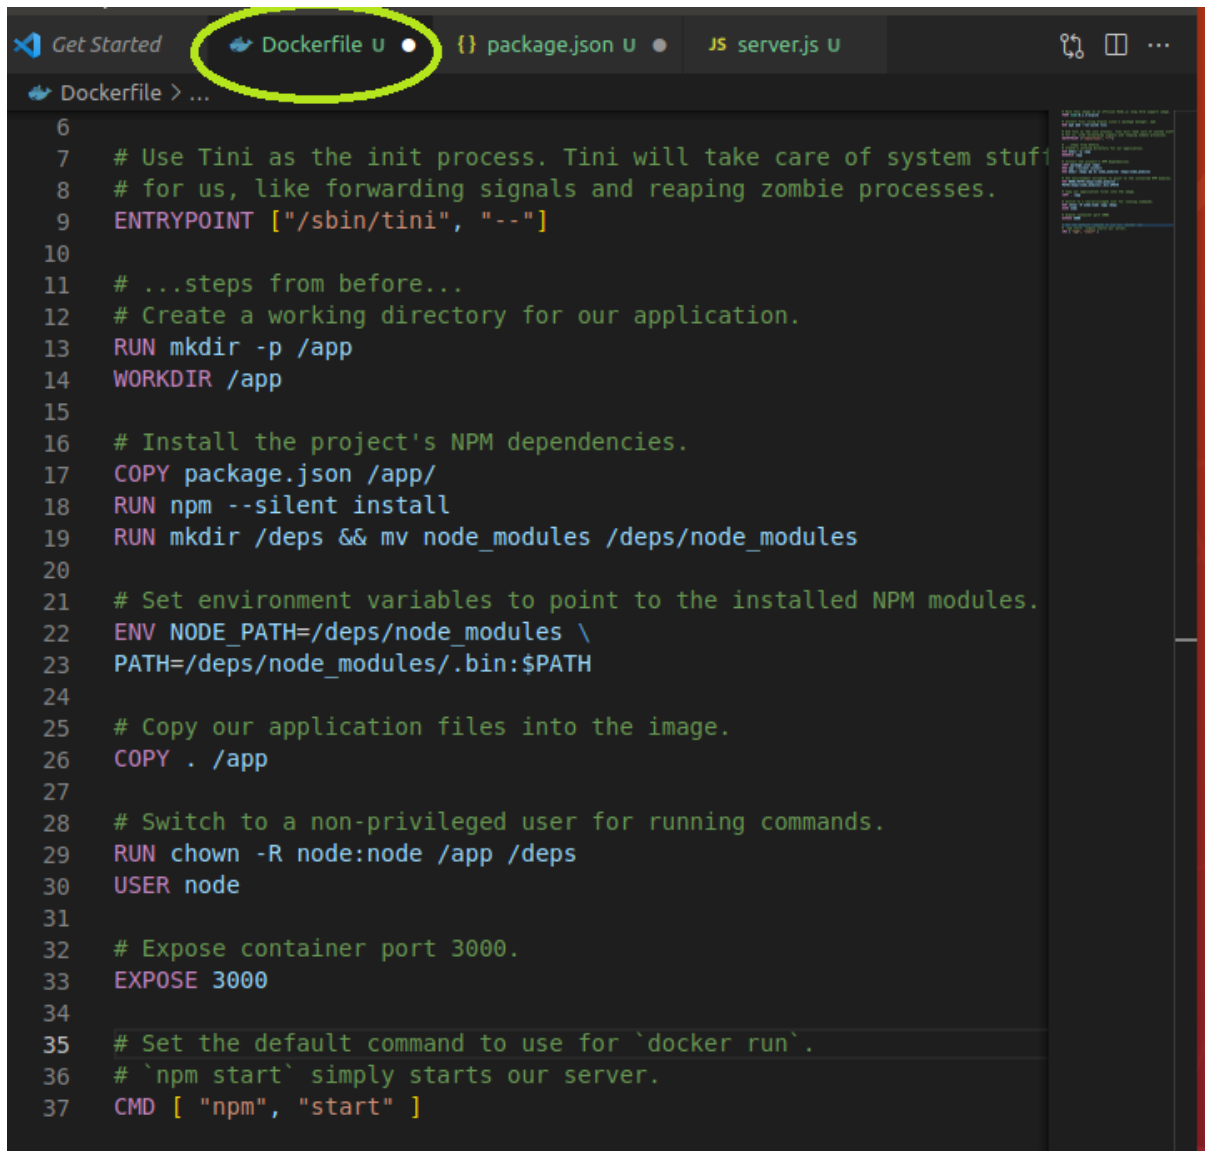
# Set environment variables to point to the installed NPM modules.
ENV NODE_PATH=/deps/node_modules \
PATH=/deps/node_modules/.bin:$PATH

# Copy our application files into the image.
COPY . /app

# Switch to a non-privileged user for running commands.
RUN chown -R node:node /app /deps
USER node

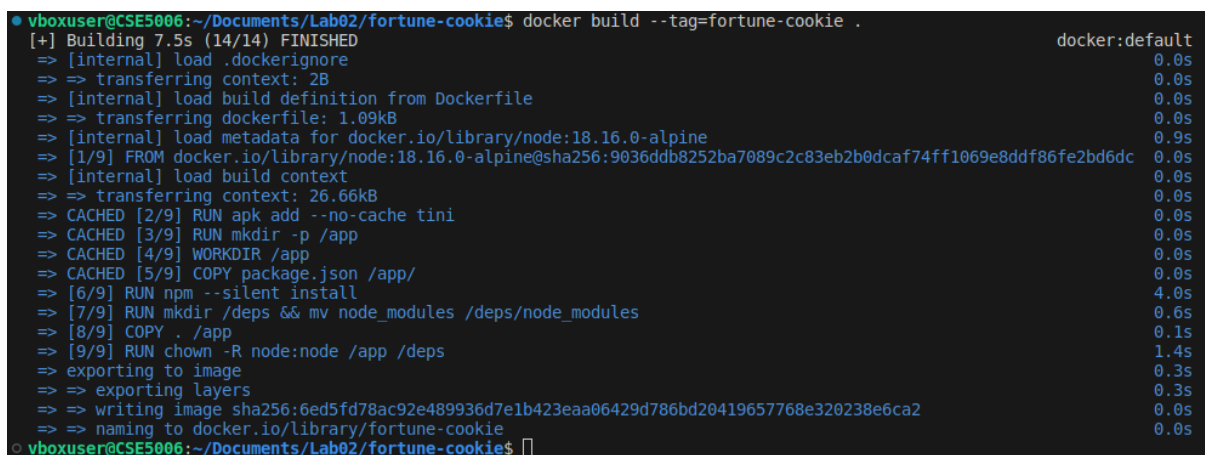
# Expose container port 3000.
EXPOSE 3000

# Set the default command to use for `docker run`.
# `npm start` simply starts our server.
CMD [ "npm", "start" ]
```



```
6
7 # Use Tini as the init process. Tini will take care of system stuff
8 # for us, like forwarding signals and reaping zombie processes.
9 ENTRYPOINT ["/sbin/tini", "--"]
10
11 # ...steps from before...
12 # Create a working directory for our application.
13 RUN mkdir -p /app
14 WORKDIR /app
15
16 # Install the project's NPM dependencies.
17 COPY package.json /app/
18 RUN npm --silent install
19 RUN mkdir /deps && mv node_modules /deps/node_modules
20
21 # Set environment variables to point to the installed NPM modules.
22 ENV NODE_PATH=/deps/node_modules \
23 PATH=/deps/node_modules/.bin:$PATH
24
25 # Copy our application files into the image.
26 COPY . /app
27
28 # Switch to a non-privileged user for running commands.
29 RUN chown -R node:node /app /deps
30 USER node
31
32 # Expose container port 3000.
33 EXPOSE 3000
34
35 # Set the default command to use for `docker run`.
36 # `npm start` simply starts our server.
37 CMD [ "npm", "start" ]
```


Build the image again, then run the container with port 3000 published on the host.



```
• vboxuser@CSE5006:~/Documents/Lab02/fortune-cookie$ docker build --tag=fortune-cookie .
[+] Building 7.5s (14/14) FINISHED
=> [internal] load .dockerignore 0.0s
=> => transferring context: 2B 0.0s
=> [internal] load build definition from Dockerfile 0.0s
=> => transferring dockerfile: 1.09kB 0.0s
=> [internal] load metadata for docker.io/library/node:18.16.0-alpine 0.9s
=> [1/9] FROM docker.io/library/node:18.16.0-alpine@sha256:9036ddb8252ba7089c2c83eb2b0dcaf74ff1069e8ddf86fe2bd6dc 0.0s
=> [internal] load build context 0.0s
=> => transferring context: 26.66kB 0.0s
=> CACHED [2/9] RUN apk add --no-cache tini 0.0s
=> CACHED [3/9] RUN mkdir -p /app 0.0s
=> CACHED [4/9] WORKDIR /app 0.0s
=> CACHED [5/9] COPY package.json /app/ 0.0s
=> [6/9] RUN npm --silent install 4.0s
=> [7/9] RUN mkdir /deps && mv node_modules /deps/node_modules 0.6s
=> [8/9] COPY . /app 0.1s
=> [9/9] RUN chown -R node:node /app /deps 1.4s
=> exporting to image 0.3s
=> => exporting layers 0.3s
=> => writing image sha256:6ed5fd78ac92e489936d7e1b423eaa06429d786bd20419657768e320238e6ca2 0.0s
=> => naming to docker.io/library/fortune-cookie 0.0s
• vboxuser@CSE5006:~/Documents/Lab02/fortune-cookie$
```

```
docker run --rm -it --publish=3000:3000 fortune-cookie
```


Publishing the port is necessary to make the host machine forward socket connections to the container image. Navigate to <http://localhost:3000> to see the application in action.



When you're done, press Ctrl+C in the terminal to terminate the web server.

Note: If you encounter any error when building the container, make sure that you have typed all the scripts correctly. Remember that copy-paste might not always work as the PDF or Words might have different encoding, especially for the non-alphanumeric symbols.

3.3.1. EXERCISE 3

Look up the Docker run reference to find out how to map container port 3000 to the default HTTP port 80 on the host. Note that **-p** is simply a shorthand for typing **--publish**. If done correctly, you should be able to access your web application in the web browser without specifying a port number.

3.4. ADDING FORTUNE COOKIE MESSAGES

Since the Docker image we're using is Alpine-based, we can add new software to the system using Alpine's package manager, **apk**. Fortunately for us (pun intended), there is a program available called **fortune**, which prints out messages in the style of a fortune cookie. We can install that program into our image by adding another **RUN** step to our Dockerfile. Insert the following **RUN** step just after the **ENTRYPOINT** step.

```
# ...  
  
# Install Fortune.  
RUN apk add --no-cache fortune  
  
# ...
```

Rebuild the image. Note that Tini is not downloaded again. This is because each step of the Dockerfile results in its own cached layer, which can be reused in subsequent builds.

Now we can tweak our **server.js** code to call **fortune** and respond with that message instead of a hello.

```

const express = require('express');
const child_process = require('child_process');

// Create a new Express web application
const app = express();

// Set handler for the index of the website
app.get('/', (req, res) => {
  // Run the system `fortune` command and respond with the message
  child_process.exec('fortune', (error, message) => {
    if(error === null) {
      res.send(message);
    } else {
      res.send('Error: ' + error);
    }
  })
});

// Start web application server
app.listen(3000, () => {
  console.log('Server started');
});

```

Build the image once more and run it. Open the website in a browser window and refresh a few times. You should be presented with a new fortune each time. Make sure that you add, commit, and push all changes to your repository at this point.

3.4.1. EXERCISE 4

Modify the server application so that it returns the current date and time along with the fortune message. In JavaScript, you can get the current date and time as a string by calling the **Date()** function. Note: you should only ever call **res.send()** once per request, so make sure that you build the entire response string and send it with a single function call.

3.5. USING DOCKER COMPOSE

Docker Compose is a handy tool that allows you to specify Docker command line options in a configuration file. This means that instead of copy-pasting a bunch of options all the time, you can specify them in a file once.

Docker Compose automatically looks for a file called **docker-compose.yml** for all of the information it needs to start your application. This file can contain one or more services, each of which results in a single container being created. Since our application only requires a single container for the moment, we will create a single entry. We will call our service "web" and nest all options under that section (**correct indentation is important**).

```
version: "3.8"

services:
  web:
    build: .
    ports:
      - "3000:3000"
```

This makes the commands for building and running much simpler. From now on, we will be using these Docker Compose commands instead of **docker build** and **docker run**.

```
docker compose build

docker compose up
```

It is not necessary to run **docker-compose build** again until the image needs to change. Besides **build** and **up**, Docker Compose has another useful command called **run**. This will start a new service instance and run the command specified. For example, if we wanted to run **fortune** directly, we could do so.

```
docker compose run --rm web fortune
```

This command will run the command **fortune** inside a web container, replacing the **CMD** line defined in the Dockerfile. The **--rm** option removes the container once it exits.

3.6. SHARING DIRECTORIES VIA VOLUMES

Currently, we copy our application source files into the image at build time using a **COPY** step in the Dockerfile. This is great for creating a production image, but it's not very convenient during development time when changes are being made rapidly. Fortunately, there is an alternative - volumes. Volumes in Docker are like shared folders. You simply specify a location on the host to map to a location in the container, and they will automatically be kept in sync while the container is running.

Add a **volumes** section to the **web** service in your **docker-compose.yml**.

```
version: "3.8"

services:
  web:
    build: .
    ports:
      - "3000:3000"
    volumes:
      - "./app"
```

This new section can be read as: "create a volume that maps the current directory on the host to the /app directory in the container."

Here's a step-by-step summary of what happens with our application files:

- During image build time,
 - Dependencies are installed in **/app/node_modules** by the **npm** install RUN step.
 - The installed dependencies are moved to **/deps/node_modules**, and the **NODE_PATH** environment variable is set appropriately so that our application knows where to find them.
 - The rest of the source files are copied into /app. This means that a fixed snapshot of the application is present inside the image for when we want to deploy the image onto a production server.
- During container start time,
 - A volume is created which maps to **/app** in the container, overwriting its previous contents with the directory on the host. We only want to do this during development, because source files are rapidly changing – in production we would just use the version of the application bundled in the image.
 - Since we moved the project dependencies to **/deps/node_modules**, our dependencies from build time will still be presented in the container.

Use Docker Compose to build the image and bring the service up. It should run like sunshine and rainbows, just as before. However, we haven't quite achieved our goal yet - try making an edit to `server.js` and reload the web page. Nothing changes... it didn't work... panic mode... PANIC! Whoa, calm down there. The changes are actually present in the container, but since we haven't restarted the Node server, it will continue using the old code already loaded in memory. To get around this limitation, we will use a wrapper program called Nodemon. Nodemon monitors source files and restarts Node when it detects a change.

1. Add "nodemon" version 1.11.0 as a dependency below Express in **package.json**, like so:

```
"dependencies": {  
  "express": "4.15.3",  
  "nodemon": "1.11.0"  
}
```

2. In the Dockerfile, change the CMD step to start Nodemon.

```
CMD [ "nodemon", "-L", "-x", "npm start" ]
```

3. Once again, build and bring the application up using Docker Compose. The output should indicate that Nodemon has started and is watching *.* (which basically means all files).
4. Make a change in **server.js**, such as adding a prefix to the message saying "Your fortune is: ". Save the file and refresh the web page in your browser (without rebuilding the image or performing any other actions). If everything went well, the changes should be immediately visible.

One thing to note is that the current setup still requires an image rebuild when new dependencies are added because downloading dependencies is part of the image creation process.

Add, commit and push your changes to GitHub once again.

3.6.1. EXERCISE 5

You may have noticed that the time shown from `Date()` is in UTC, not local time. This is because the default time zone set in the `node:18.16.0-alpine` Docker image is UTC. One way to change this behavior is to mount the time zone settings from the host into the Docker image. Add a line to `docker-compose.yml` which mounts `/etc/localtime` into the container.

```
volumes:
  - "./app"
  - "/etc/localtime:/etc/localtime:ro"
```

The "ro" part tells Docker that we want read-only access to the file - this prevents the container from being able to change the host's time zone setting. Rebuild and run the application. Refresh the web page and check that the time is displayed in local time (you should see AEST in parentheses).